

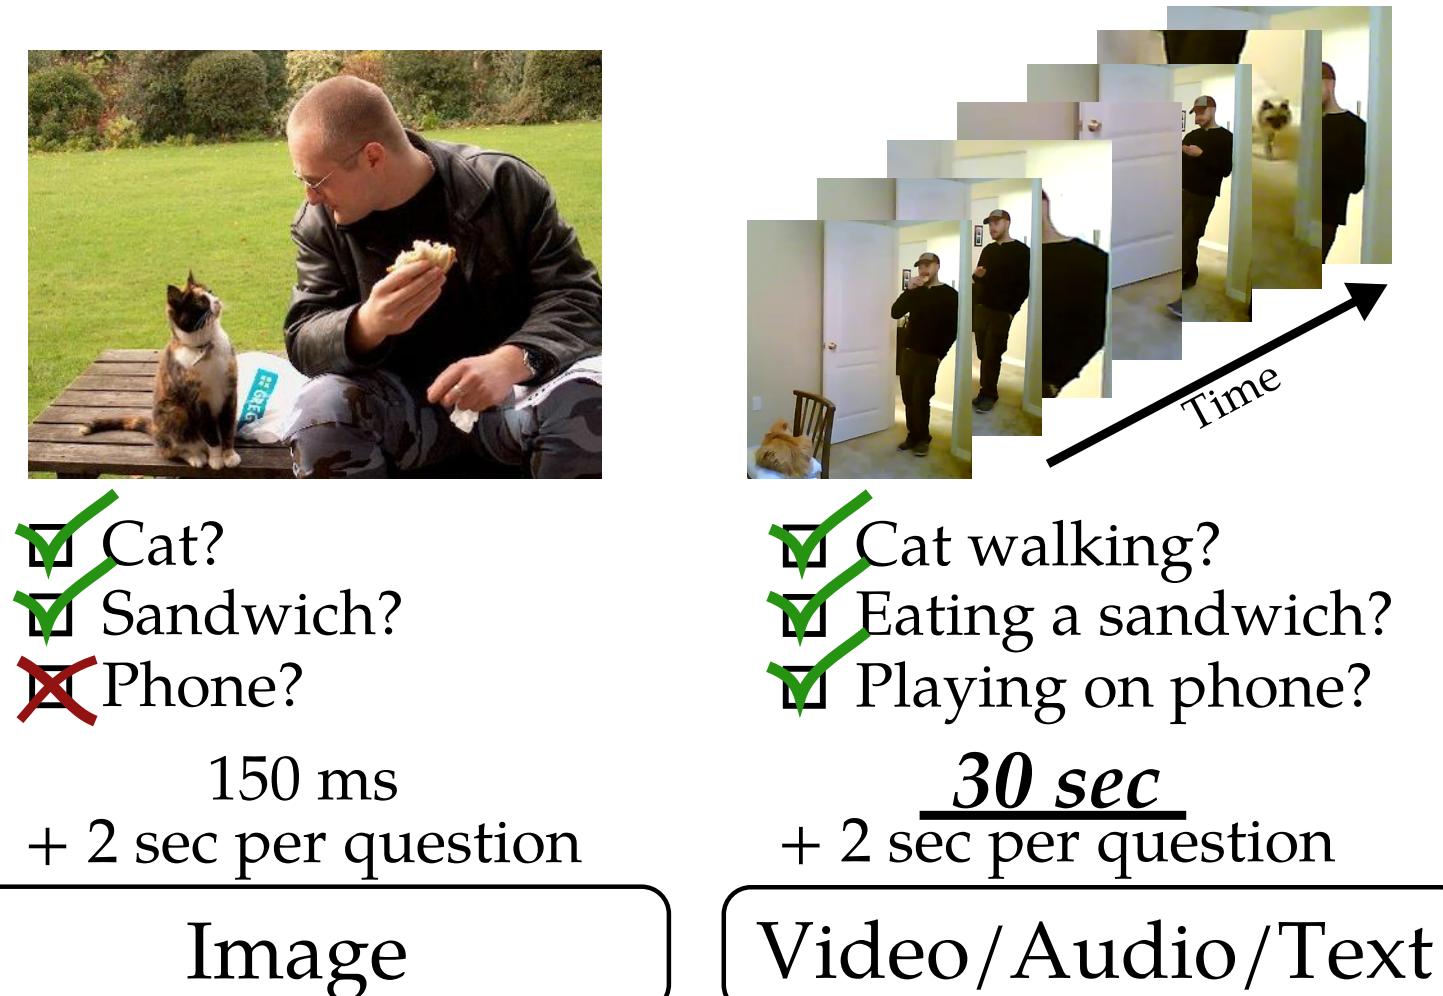
Much Ado About Time: Exhaustive Annotation of Temporal Data

Gunnar A. Sigurdsson (CMU), Olga Russakovsky (CMU), Ali Farhadi (UW, AI2), Ivan Laptev (INRIA), Abhinav Gupta (CMU)

<http://allenai.org/plato/charades/>

Published in HCOMP'16

Goal: explore cost-effective strategies for annotating temporal data (e.g., videos)



How many questions about a video can we ask simultaneously?

Few questions: questions can be kept in worker's short-term memory, increasing annotation recall

Many questions: video watching time can be amortized, reducing annotation cost

Instructions

Below is a link to a video of one or two people, please watch each video and answer the questions.

- This HIT contains multiple videos, each followed by few questions. The number of videos and questions is balanced such that the task should take 3 minutes.
- Make sure you **fully and carefully watch each video** so you do not miss anything. **This is important**.
- It is possible that many of the actions in this HIT do not match. It is important to verify an action is indeed **not** present in the video.
- Check all that apply! If there is any doubt, check it anyway for good measure.
- Read each and every question carefully. Do not take shortcuts, it will cause you to miss something.

Check here if someone is **Taking a picture of something** in the video
 Check here if someone is **interacting with cup/glass/bottle** in the video
If checked, how? (Select all that apply. Use ctrl or cmd to select multiple):
Drinking from a cup/glass/bottle
Holding a cup/glass/bottle of something
Pouring something into a cup/glass/bottle
Putting a cup/glass/bottle somewhere
Taking a cup/glass/bottle from somewhere
Washing a cup/glass/bottle
Other

 Check here if someone is **interacting with laptop** in the video
 Check here if someone is **interacting with doorknob** in the video
 Check here if someone is **interaction with table** in the video

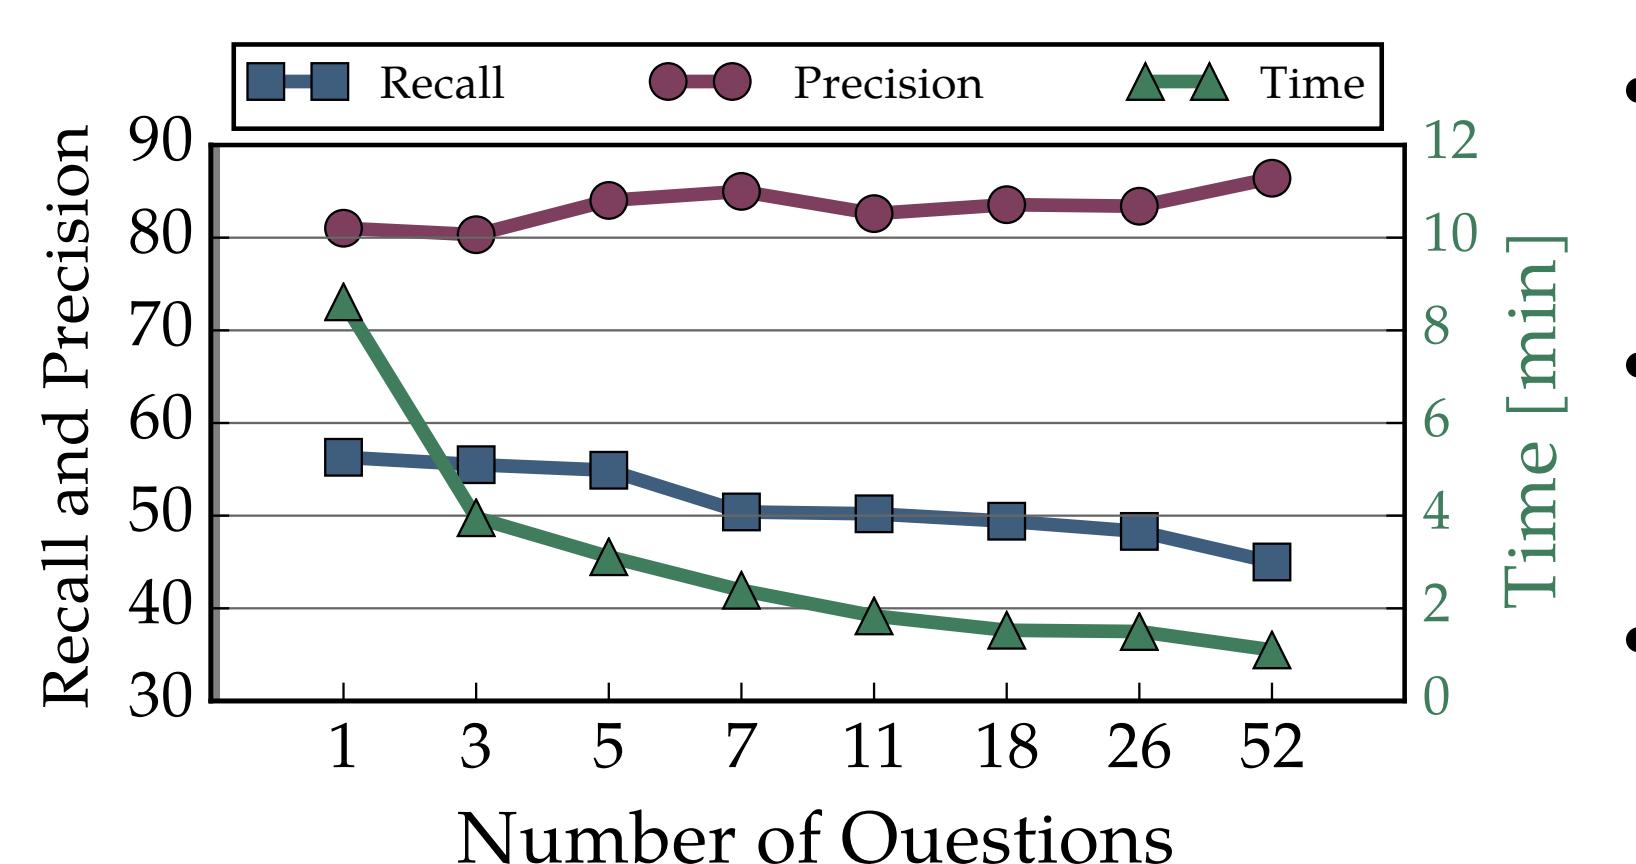
Dataset: Charades video dataset of [Sigurdsson ECCV16]



- Household activities
- Videos recorded by AMT workers
- 9,848 videos, averaging 30 secs
- 157 activity classes (52 hierarchically)
- Some activities pre-annotated but not exhaustively

Experiment #1: Run one round of annotation

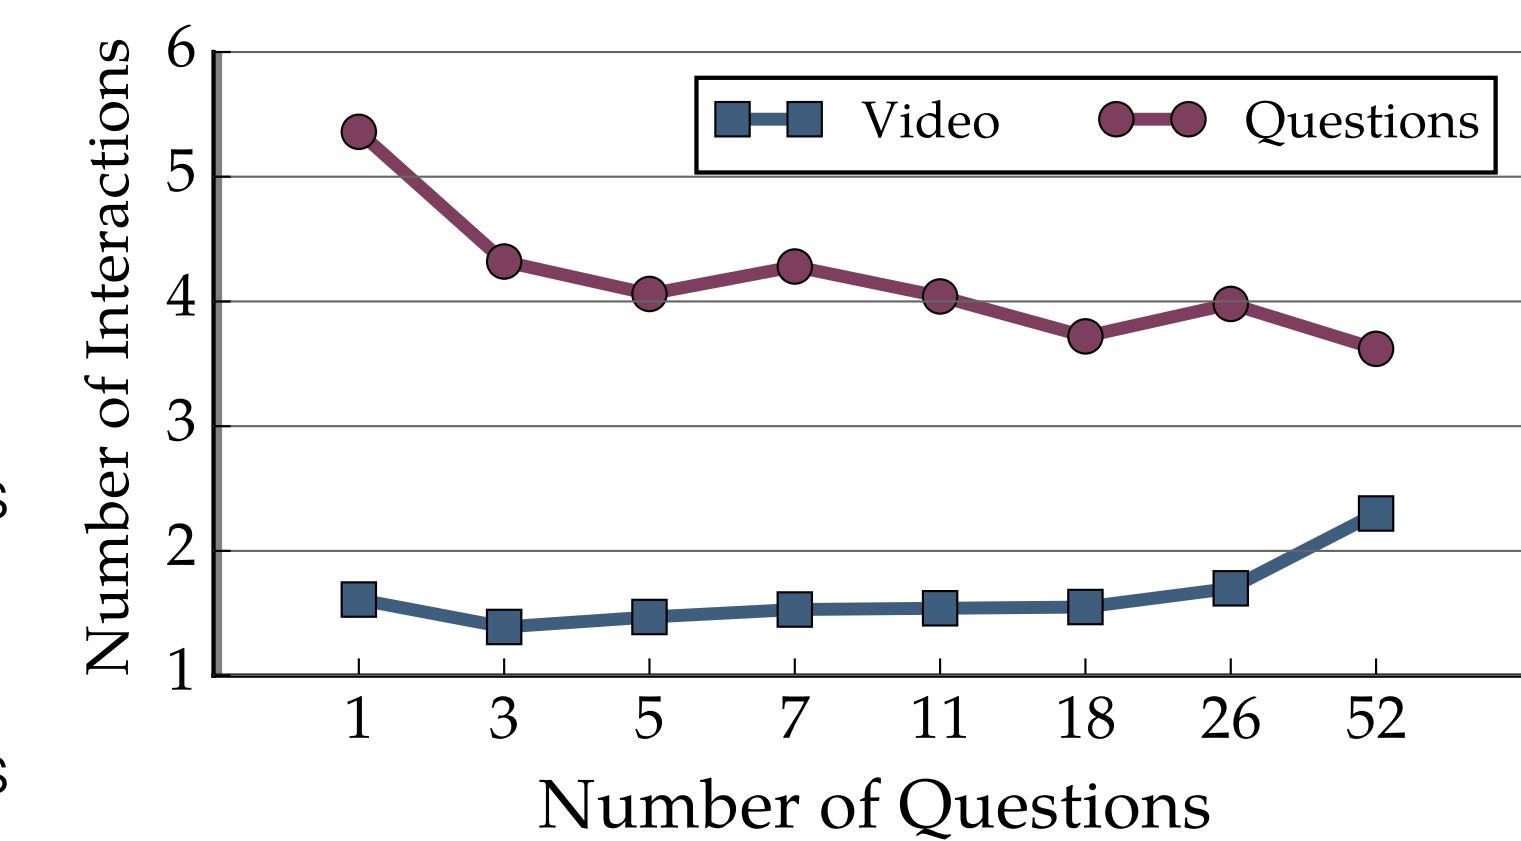
Annotate 157 actions in 140 videos on Amazon Mechanical Turk



- Recall is *lower* with more questions, as expected
- Precision turns out to be *higher* with more questions
- Time/cost is *lower* with more questions, as expected

Analyze worker behavior during annotation

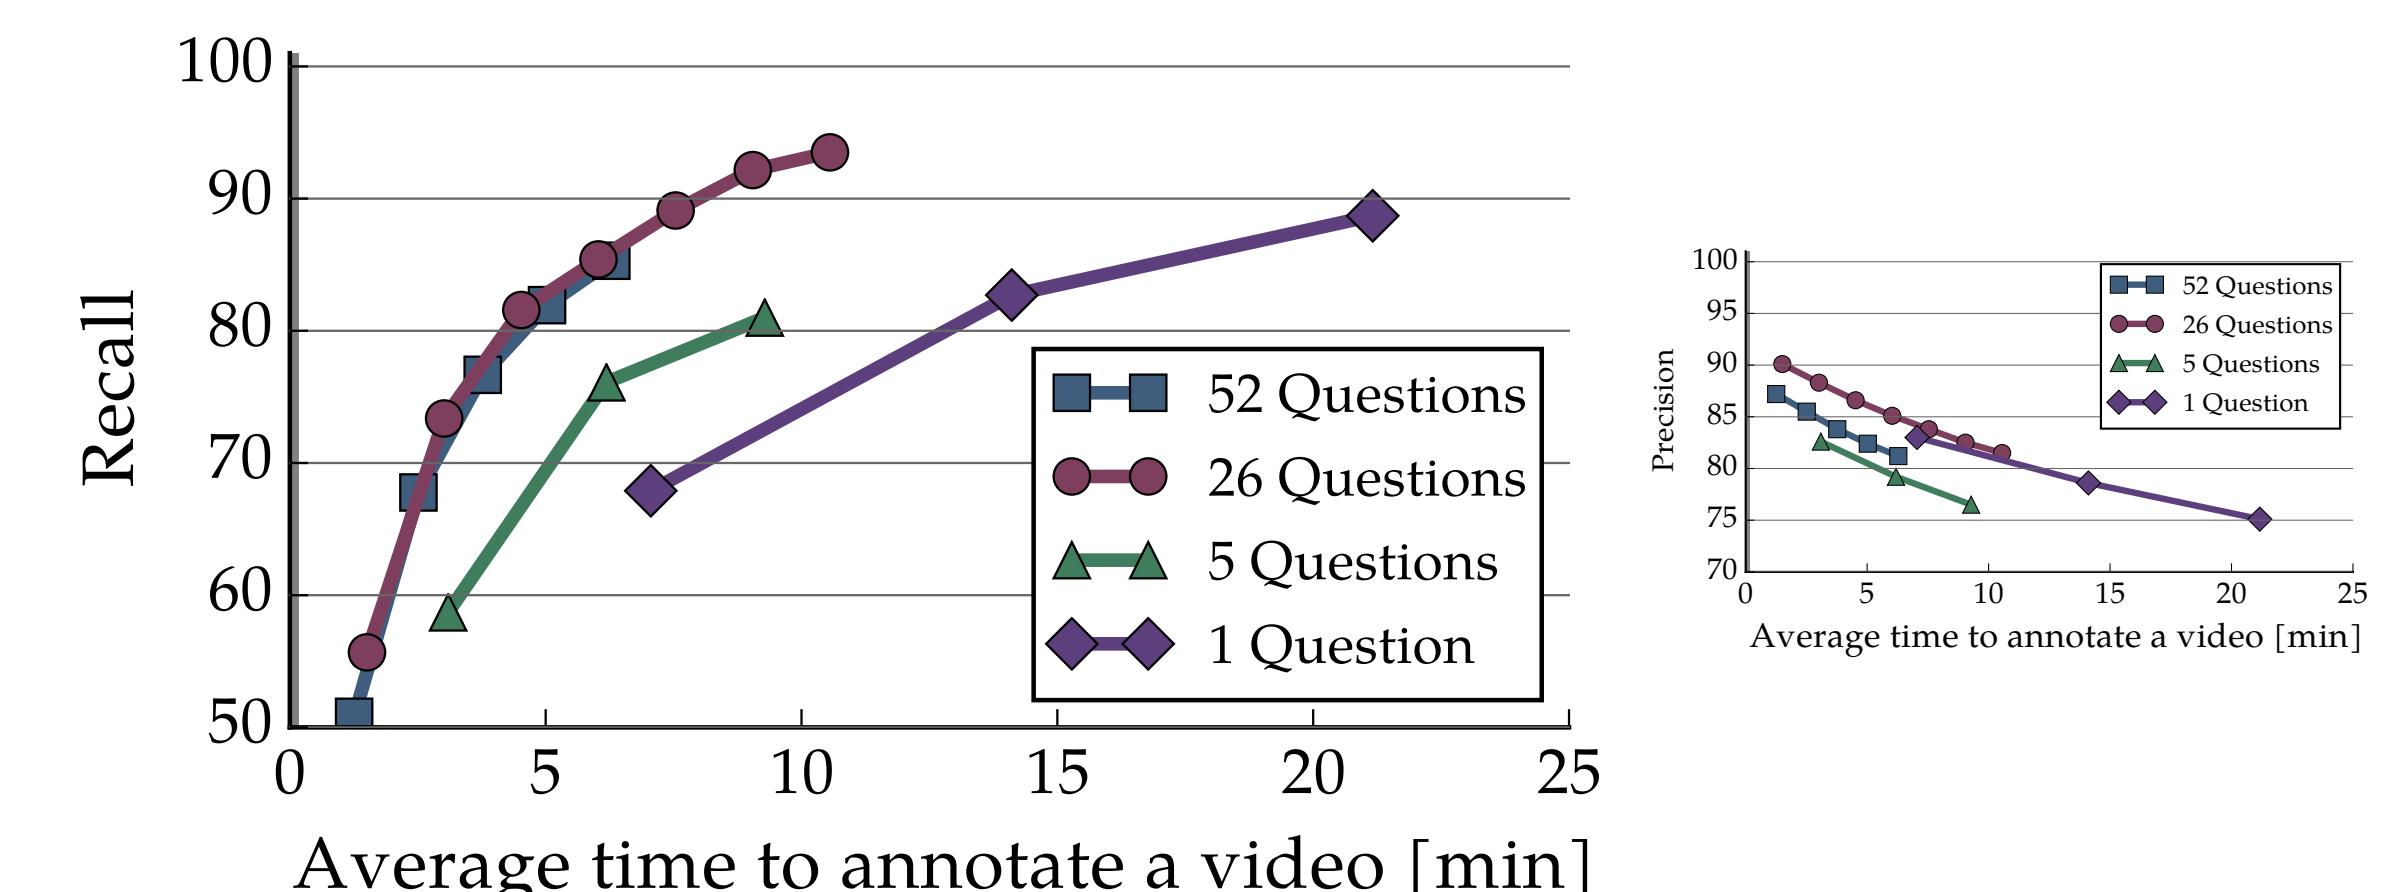
- Worker enjoyment doesn't depend on number of questions
- Workers **interact** with the video more when asked more questions
- Workers annotate fewer actions when asked more questions



Experiment #2: Improve and re-evaluate the UI

- Improving time/cost: Keep questions consistent within a HIT
 - 13.6% reduction in annotation time
- Improving precision: Use the available annotations to inject more positive questions into the HIT
 - 2.3% improvement in precision
- Improving recall try #1: First request a 20-word video summary
 - No effect on recall but 40% slower annotation
- Improving recall try #2: Force yes/no response to every question
 - Recall actually drops (annoys workers?)
- Improving recall for real: Run multiple rounds of annotation
 - Recall improves from 58.0% to 83.3% after 3 rounds

Experiment #3: Bringing it all together



Conclusion: Asking **many questions** per video is a much more effective use of worker time than asking just a **few questions**

Annotated data: 9,848 videos exhaustively annotated

- 66,963 action instances from 1,310,014 collected annotations
- Increased density of annotation from 3.7 to 9.0 labels per video
- Additionally collected temporal annotations
- All publicly available at <http://allenai.org/plato/charades/>

This work was supported by ONR MURI N00014-16-1-2007, ERC award ACTIVIA, Allen Distinguished Investigator Award, and the Allen Institute for Artificial Intelligence.